

# Prácticas BigData

## 1. Comando YARN

- Vamos a ver como trabajar con el comando YARN
- Mantenemos dos terminales abiertos
- En el primero lanzamos un proyecto mapreduce.
  - Si es en versiones anteriores a la 2.9 tenemos que lanzarlo de esta manera. Es debido a que en la 2.9 se ha modificado la propiedad queue.name por queuename

```
hadoop jar /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.9.0.jar wordcount -Dmapred.job.queue.name=warehouse /practicascite75_99.txt /salida20
```

- En la 2.9 lo hacemos así

```
hadoop jar /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.9.0.jar wordcount -Dmapred.job.queue.name=warehouse /practicascite75_99.txt /salida20
```

En el segundo vemos la aplicación con el comando yarn

- Vamos a comprobar que se ha lanzado

**yarn application -list**

18/01/07 20:57:28 INFO client.RMPProxy: Connecting to ResourceManager at nodo1/192.168.56.101:8032

Total number of applications (application-types: [], states: [SUBMITTED, ACCEPTED, RUNNING] and tags: []):1

Queue	Application-Id State	Application-Name Final-State	Application-Type Progress	User Tracking-URL
application_1515354119376_0002	word count	MAPREDUCE		
hadoop default	ACCEPTED	UNDEFINED	0%	
N/A				

- Comprobamos en que nodos se está ejecutando

**yarn node -list**

18/01/07 20:57:38 INFO client.RMPProxy: Connecting to ResourceManager at nodo1/192.168.56.101:8032

Total Nodes:6

Node-Id	Node-State	Node-Http-Address	Number-of-Running-Containers
nodo7:34788	RUNNING	nodo7:8042	0
nodo3:40979	RUNNING	nodo3:8042	0
nodo4:36792	RUNNING	nodo4:8042	0
nodo6:33922	RUNNING	nodo6:8042	0
nodo2:36376	RUNNING	nodo2:8042	3
nodo5:37517	RUNNING	nodo5:8042	0

- Comprobamos el estado completo de la aplicación

**yarn application -status application\_1515354119376\_0002**

18/01/07 20:57:50 INFO client.RMPProxy: Connecting to ResourceManager at nodo1/192.168.56.101:8032

Application Report :

Application-Id : application\_1515354119376\_0002

Application-Name : word count

Application-Type : MAPREDUCE

User : hadoop

Queue : default

Application Priority : 0

Start-Time : 1515355041442

Finish-Time : 0

Progress : 5%

State : RUNNING

Final-State : UNDEFINED

Tracking-URL : http://nodo2:34944

RPC Port : 41274

AM Host : nodo2

Aggregate Resource Allocation : 91429 MB-seconds, 59 vcore-seconds

Aggregate Resource Preempted : 0 MB-seconds, 0 vcore-seconds

Log Aggregation Status : DISABLED

Diagnostics :

Unmanaged Application : false

Application Node Label Expression : <Not set>

AM container Node Label Expression : <DEFAULT\_PARTITION>

TimeoutType : LIFETIME      ExpiryTime : UNLIMITED      RemainingTime : -1seconds

- Comprobamos los Application Attempts:

**yarn applicationattempt -list application\_1515354119376\_0002**

18/01/07 20:58:50 INFO client.RMPProxy: Connecting to ResourceManager at nodo1/192.168.56.101:8032

Total number of application attempts :1

ApplicationAttempt-Id	State	AM-Container-Id
Tracking-URL		
appattempt_1515354119376_0002_000001	RUNNING	
container_1515354119376_0002_01_000001		
http://nodo1:8088/proxy/application_1515354119376_0002/		

- Vemos los contenedores asociados

**yarn container -list appattempt\_1515354119376\_0002\_000001**

```
18/01/07 20:59:13 INFO client.RMPProxy: Connecting to ResourceManager at
nodo1/192.168.56.101:8032
```

Total number of containers :2

Host	Container-Id	Start Time	Finish Time	State
Node	Http Address		LOG-URL	
container_1515354119376_0002_01_000004	dom ene 07 20:58:51 +0100 2018			
N/A	RUNNING	nodo2:36376	http://nodo2:8042	
	http://nodo2:8042/node/containerlogs/container_1515354119376_0002_01_000004/hadoop			
container_1515354119376_0002_01_000001	dom ene 07 20:57:21 +0100 2018			
N/A	RUNNING	nodo2:36376	http://nodo2:8042	
	http://nodo2:8042/node/containerlogs/container_1515354119376_0002_01_000001/hadoop			

- Eliminamos la aplicación

```
yarn application -kill application_1515354119376_0002
```

- Comprobamos que se ha eliminado